

Modeling Streaming Data for Processing with Apache Beam

GETTING STARTED WITH STREAM PROCESSING



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Batch data and bounded datasets

Streaming data for unbounded datasets and real-time processing

Micro-batch processing and continuous processing

Lambda and Kappa architectures

Challenges in real-time stream processing

Prerequisites and Course Outline

Prerequisites



No prior experience of working with Streaming Data required

Experience programming in Java

Apache Maven for dependency management

Course Outline



Getting Started with Stream Processing

Introducing Apache Beam for Stream Processing

Perform Windowing Operations

Batch Processing and Stream Processing

Analysis of Deliveries



E-commerce site

**How are they distributed
across the country?**

**Are there routes that can
be clubbed together?**

**How do courier
companies compare?**



Analysis of Deliveries



Generate **periodic** reports to improve delivery metrics

Analysis of Deliveries

Collect

Source and destination of packages

Courier company details

Extract

Trends in the form of visuals

Actionable insights

Analyze

Run jobs on different slices

Courier, metro areas, rural areas, warehouses

Analysis of Deliveries



Bounded datasets: Finite unchanging datasets to analyze

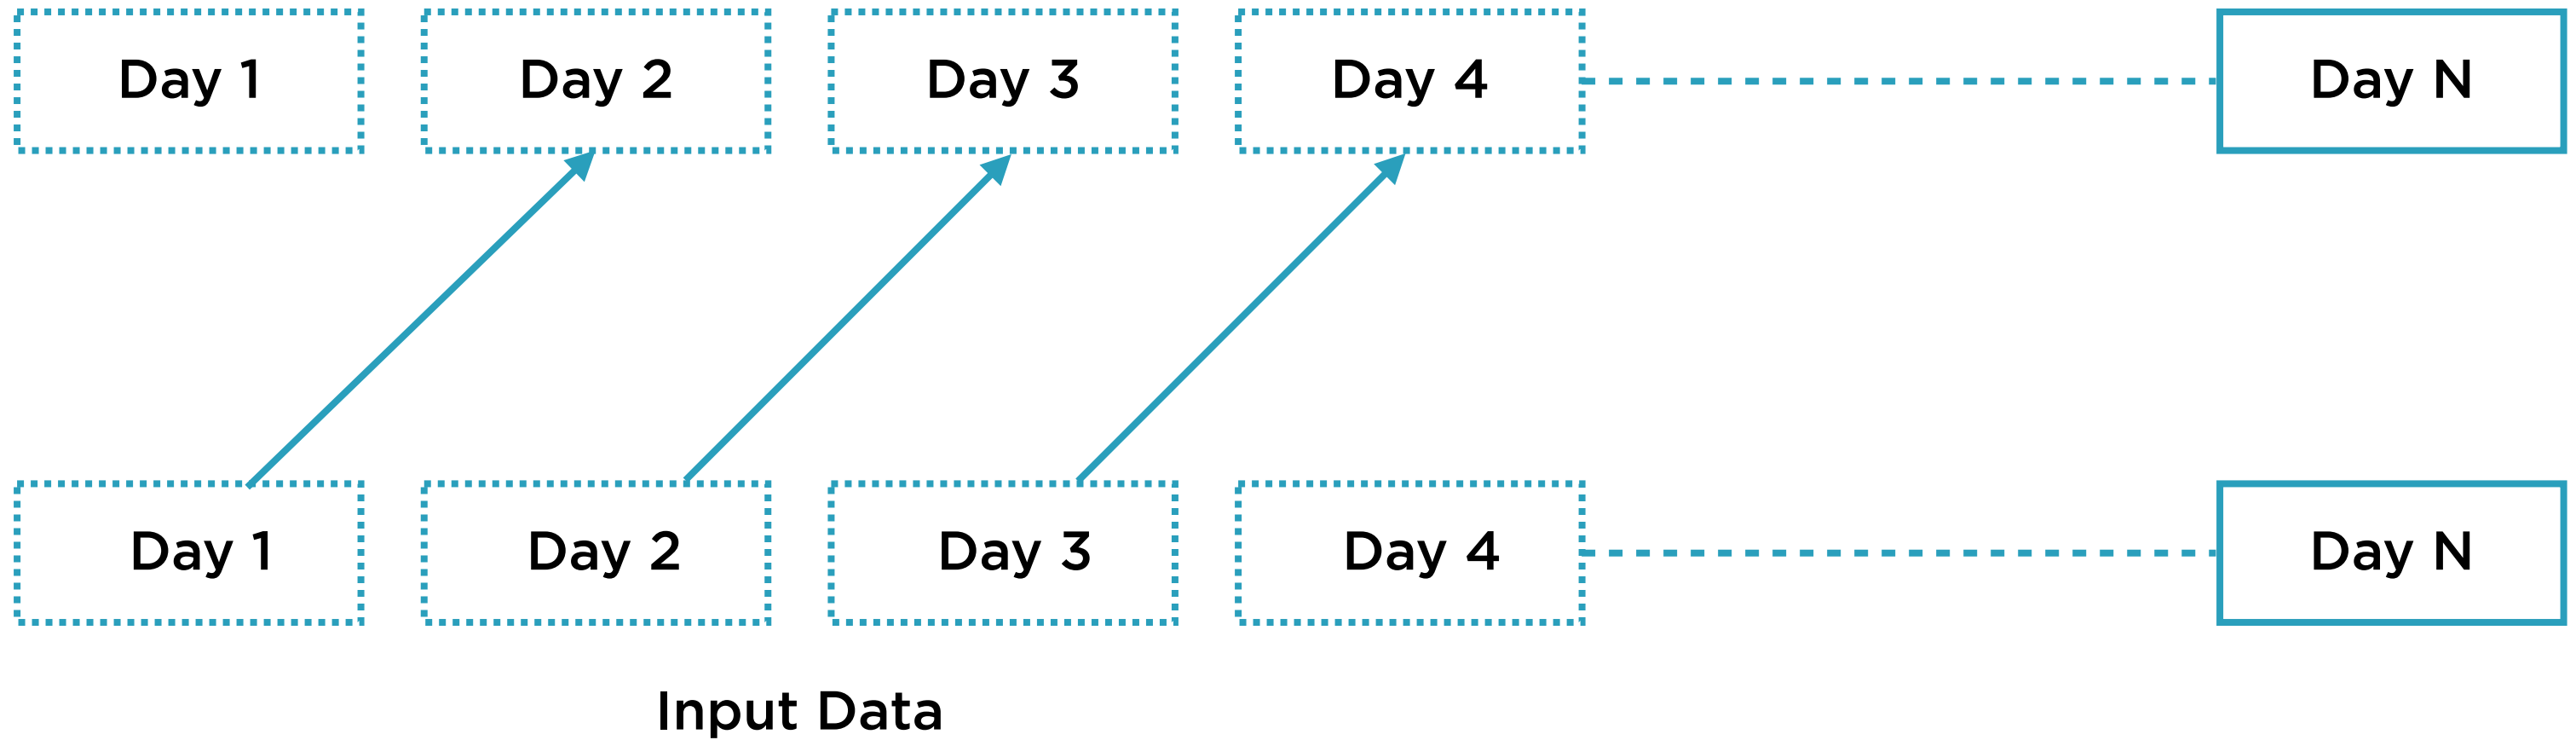
- week, month, year

Batch processing: Runs for a specific time, completes, releases resources

- minutes, hours, days

Batch Processing

Processing Data



Batch Processing

Processing Data


Day 1

Day 1

Input Data


Batch Processing

Processing Data




Input Data

**Stored data processed over a
period of time**




Batch Processing

Processing Data




Input Data

**Stored data processed over a
period of time**




Batch Processing

Processing Data



Input Data

Stored data processed over a
period of time




Batch Processing

Processing Data

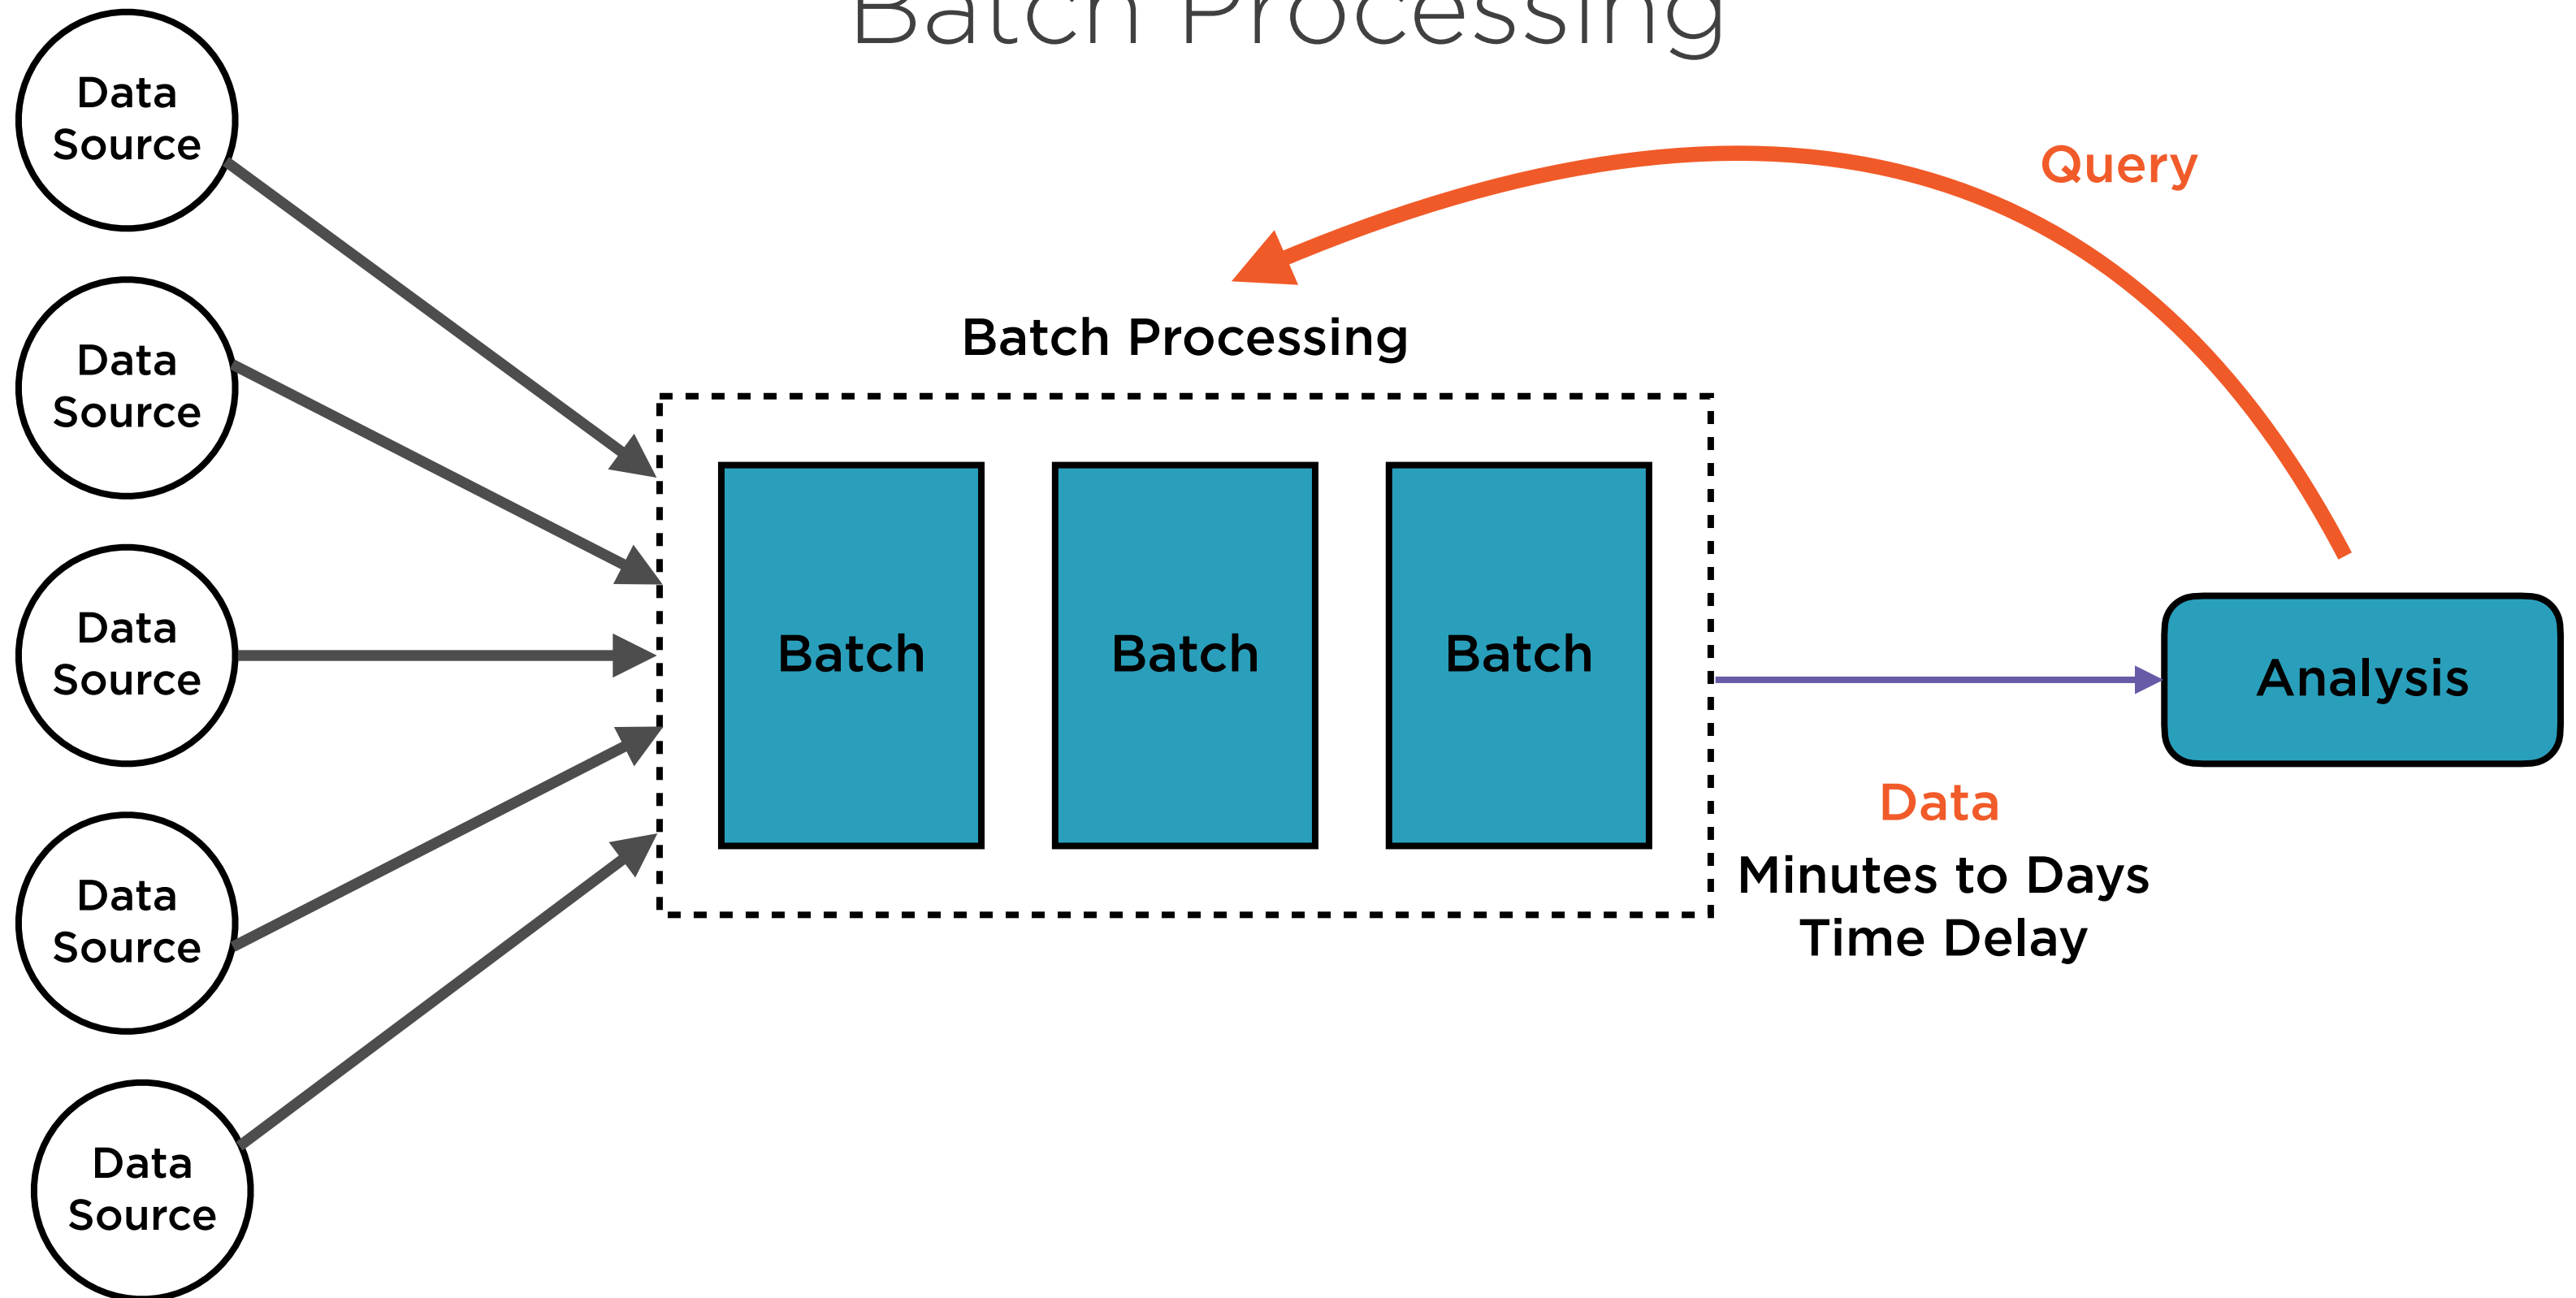


Input Data

Stored data processed over a
period of time



Batch Processing



Tracking of Deliveries



E-commerce site

**Real-time location of
delivery agents**

**Real-time order status
updates**

**Real-time inventory
tracking**



Tracking of Deliveries



Continuously monitor data to ensure deliveries are flowing smoothly

Tracking of Deliveries

Monitor

Constantly listen for updates
GPS coordinates, status
information, inventory changes

Extract

Plot real-time graphs
Track on a map

Process

As entities flow in process them
in micro-batches
Whole stream, predetermined
window



Tracking of Deliveries



Unbounded datasets: Infinite datasets which are added to continuously

- streaming data

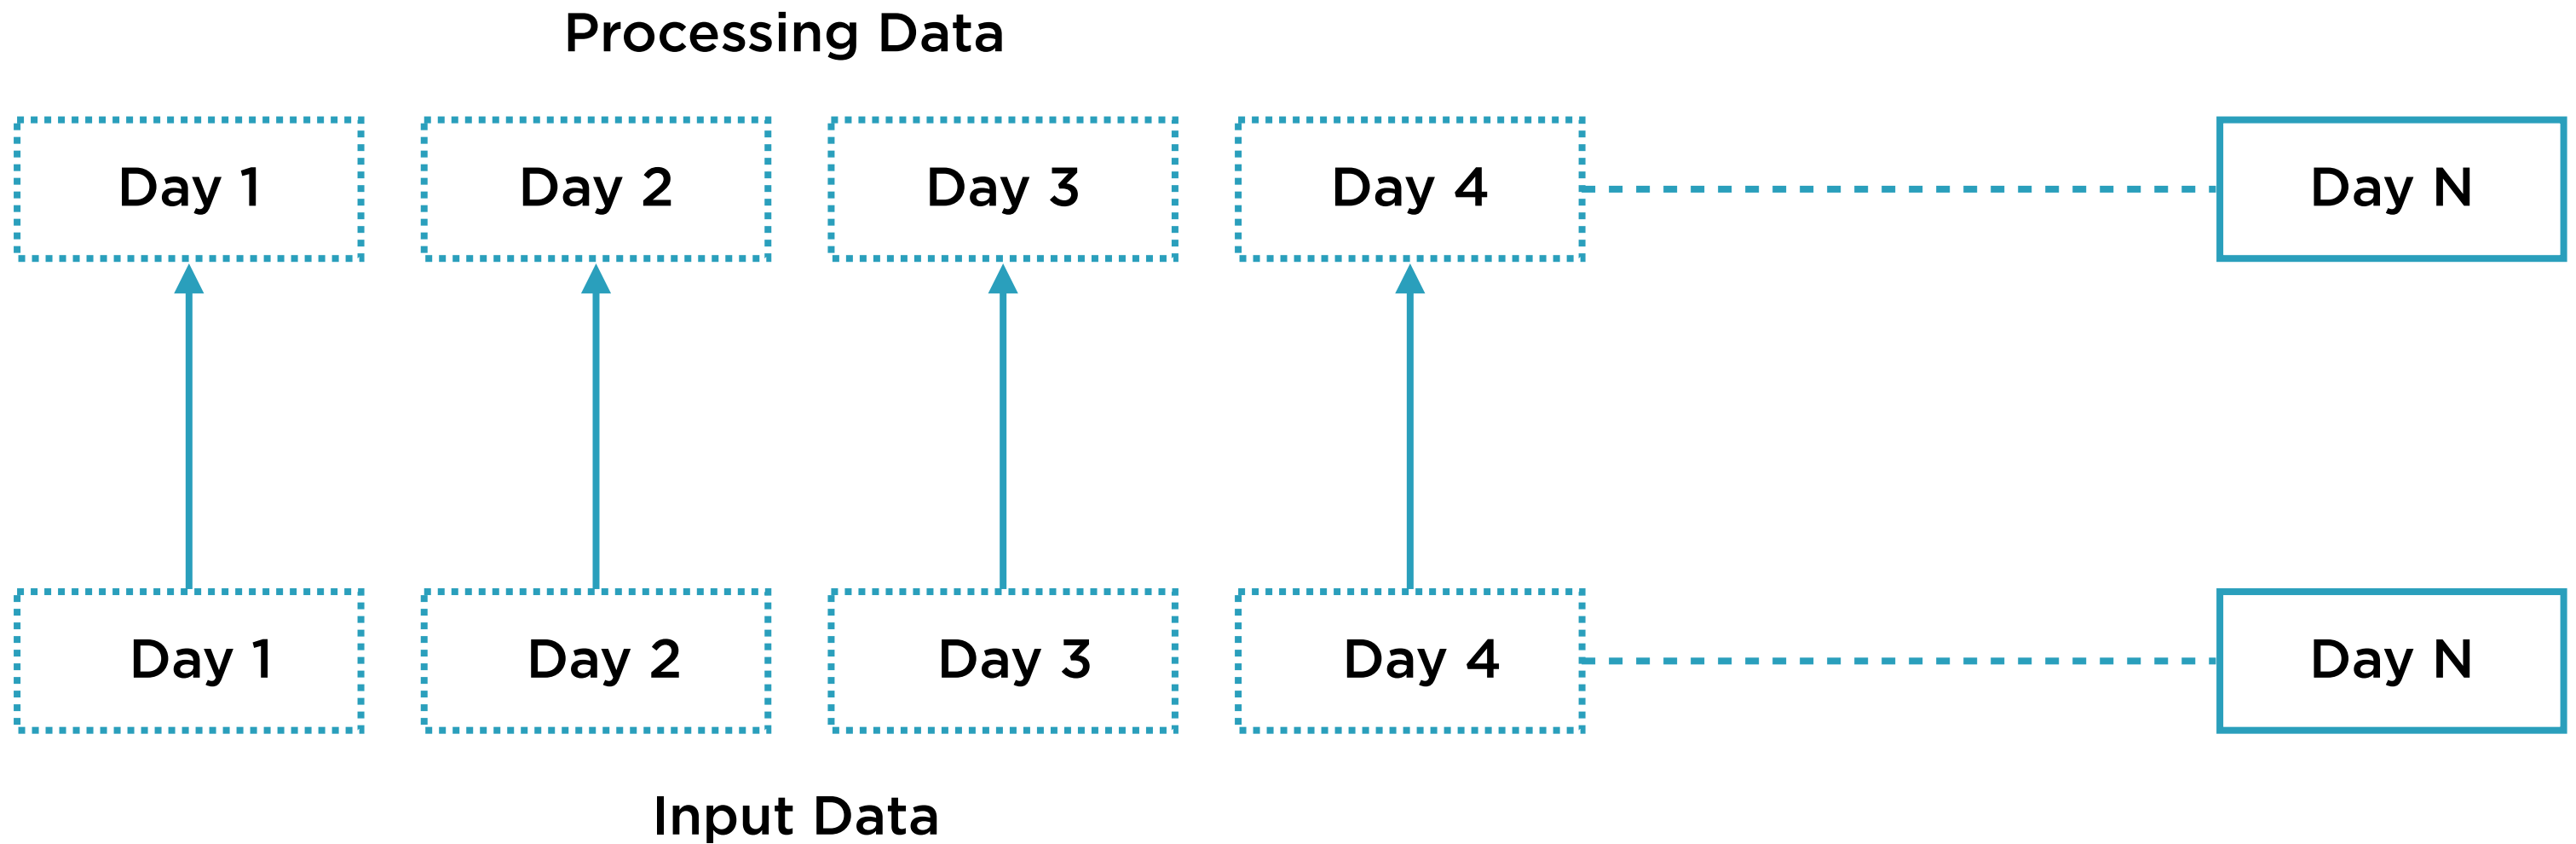
Continuous processing: Runs constantly as long as data is received

- stream processing

Bounded datasets are
processed in **batches**


Unbounded datasets are
processed as **streams**

Stream Processing



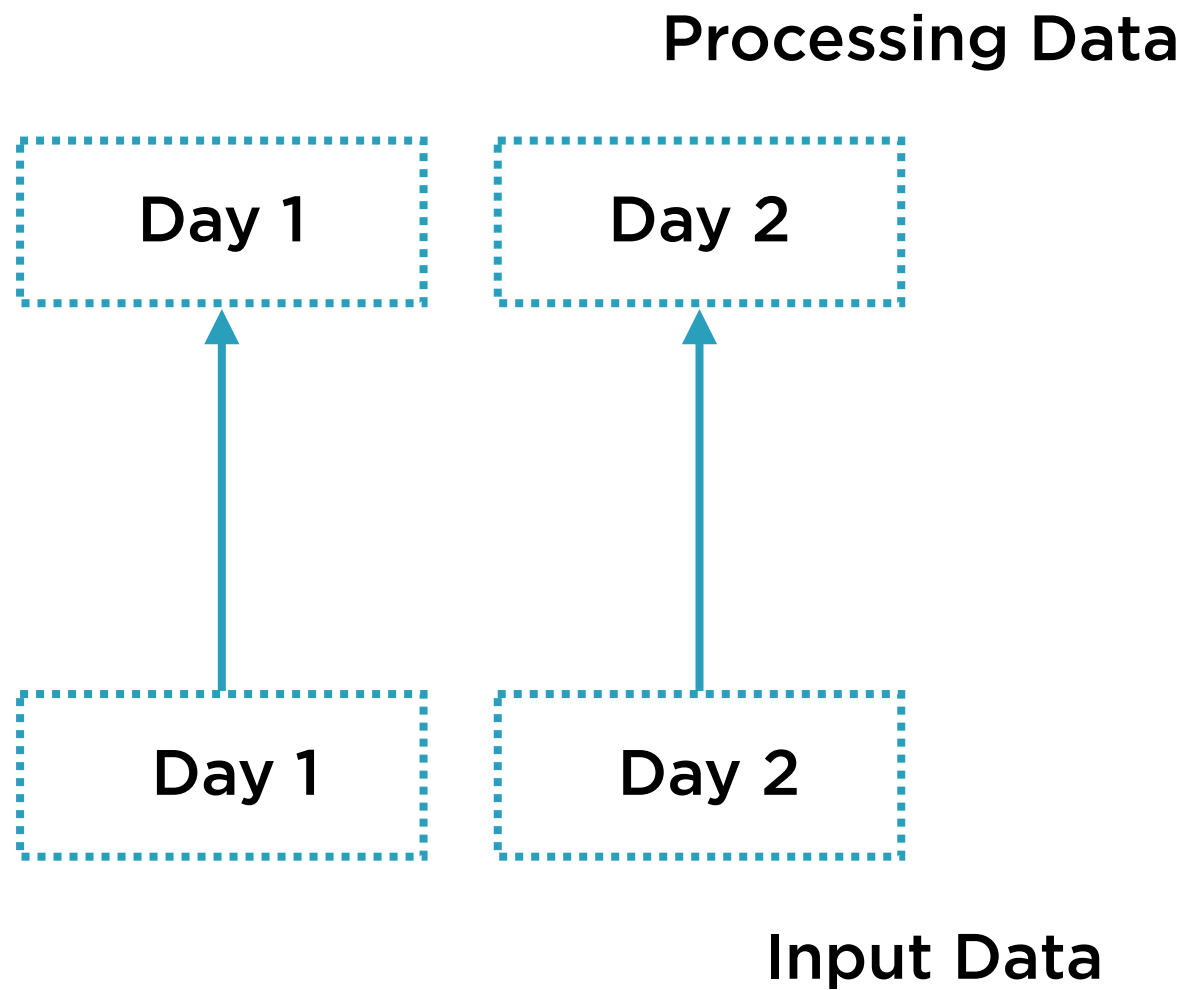
Stream Processing

Processing Data



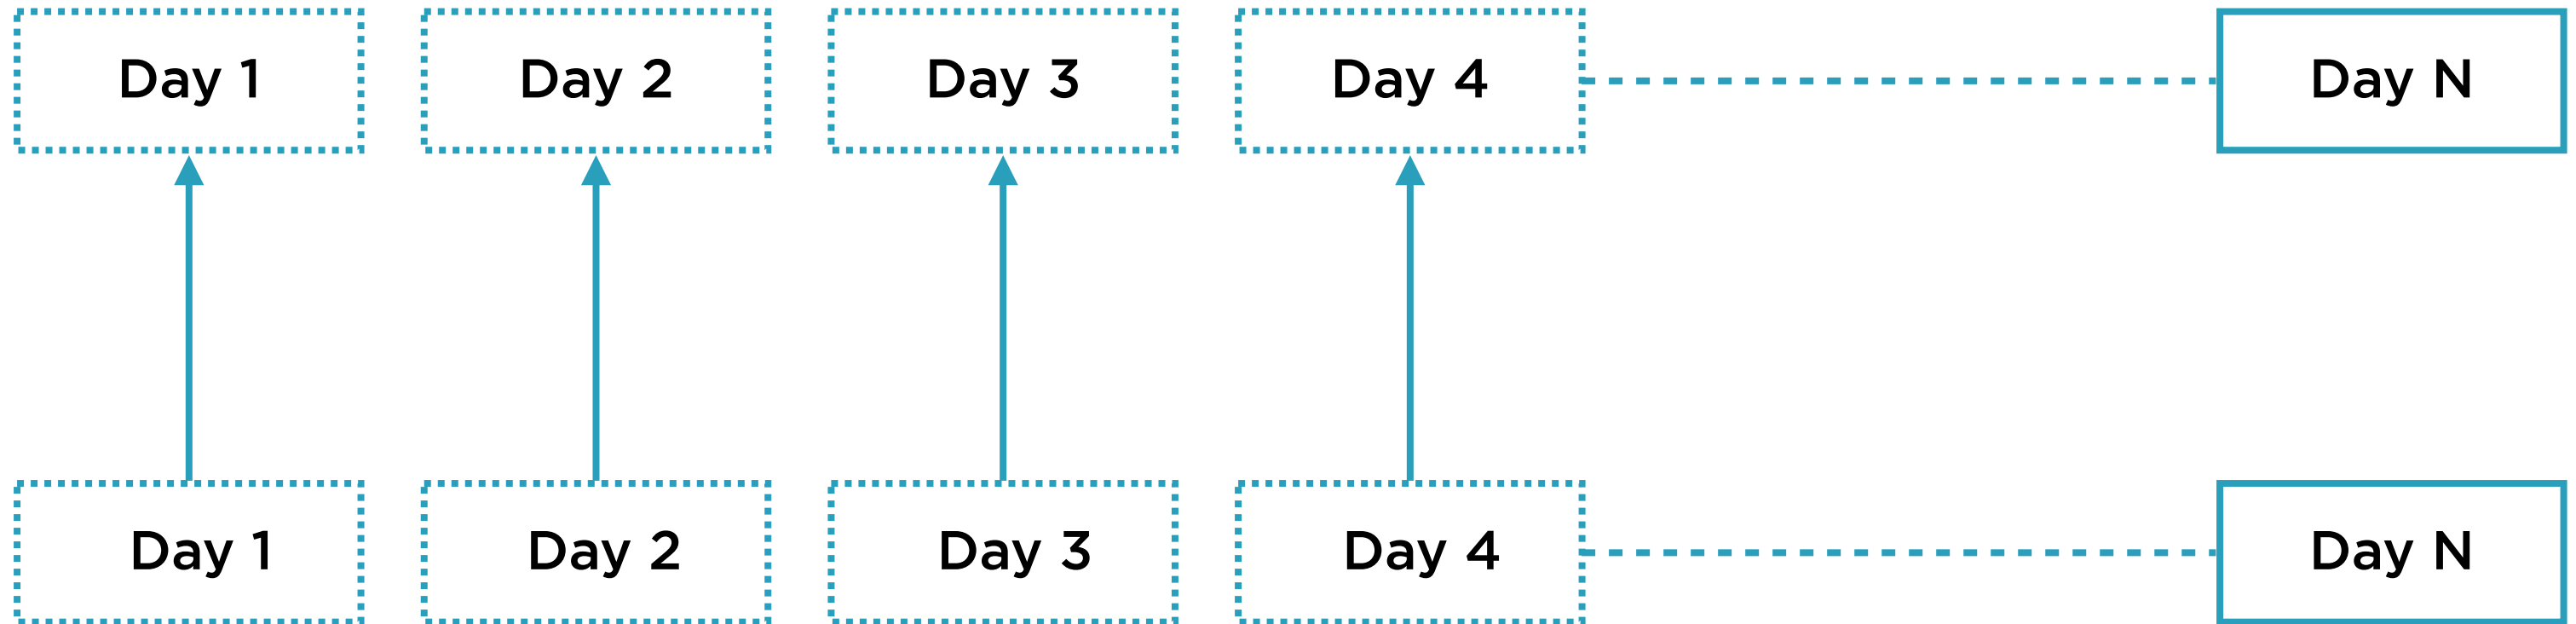
Input Data

Stream Processing




Stream Processing

Processing Data

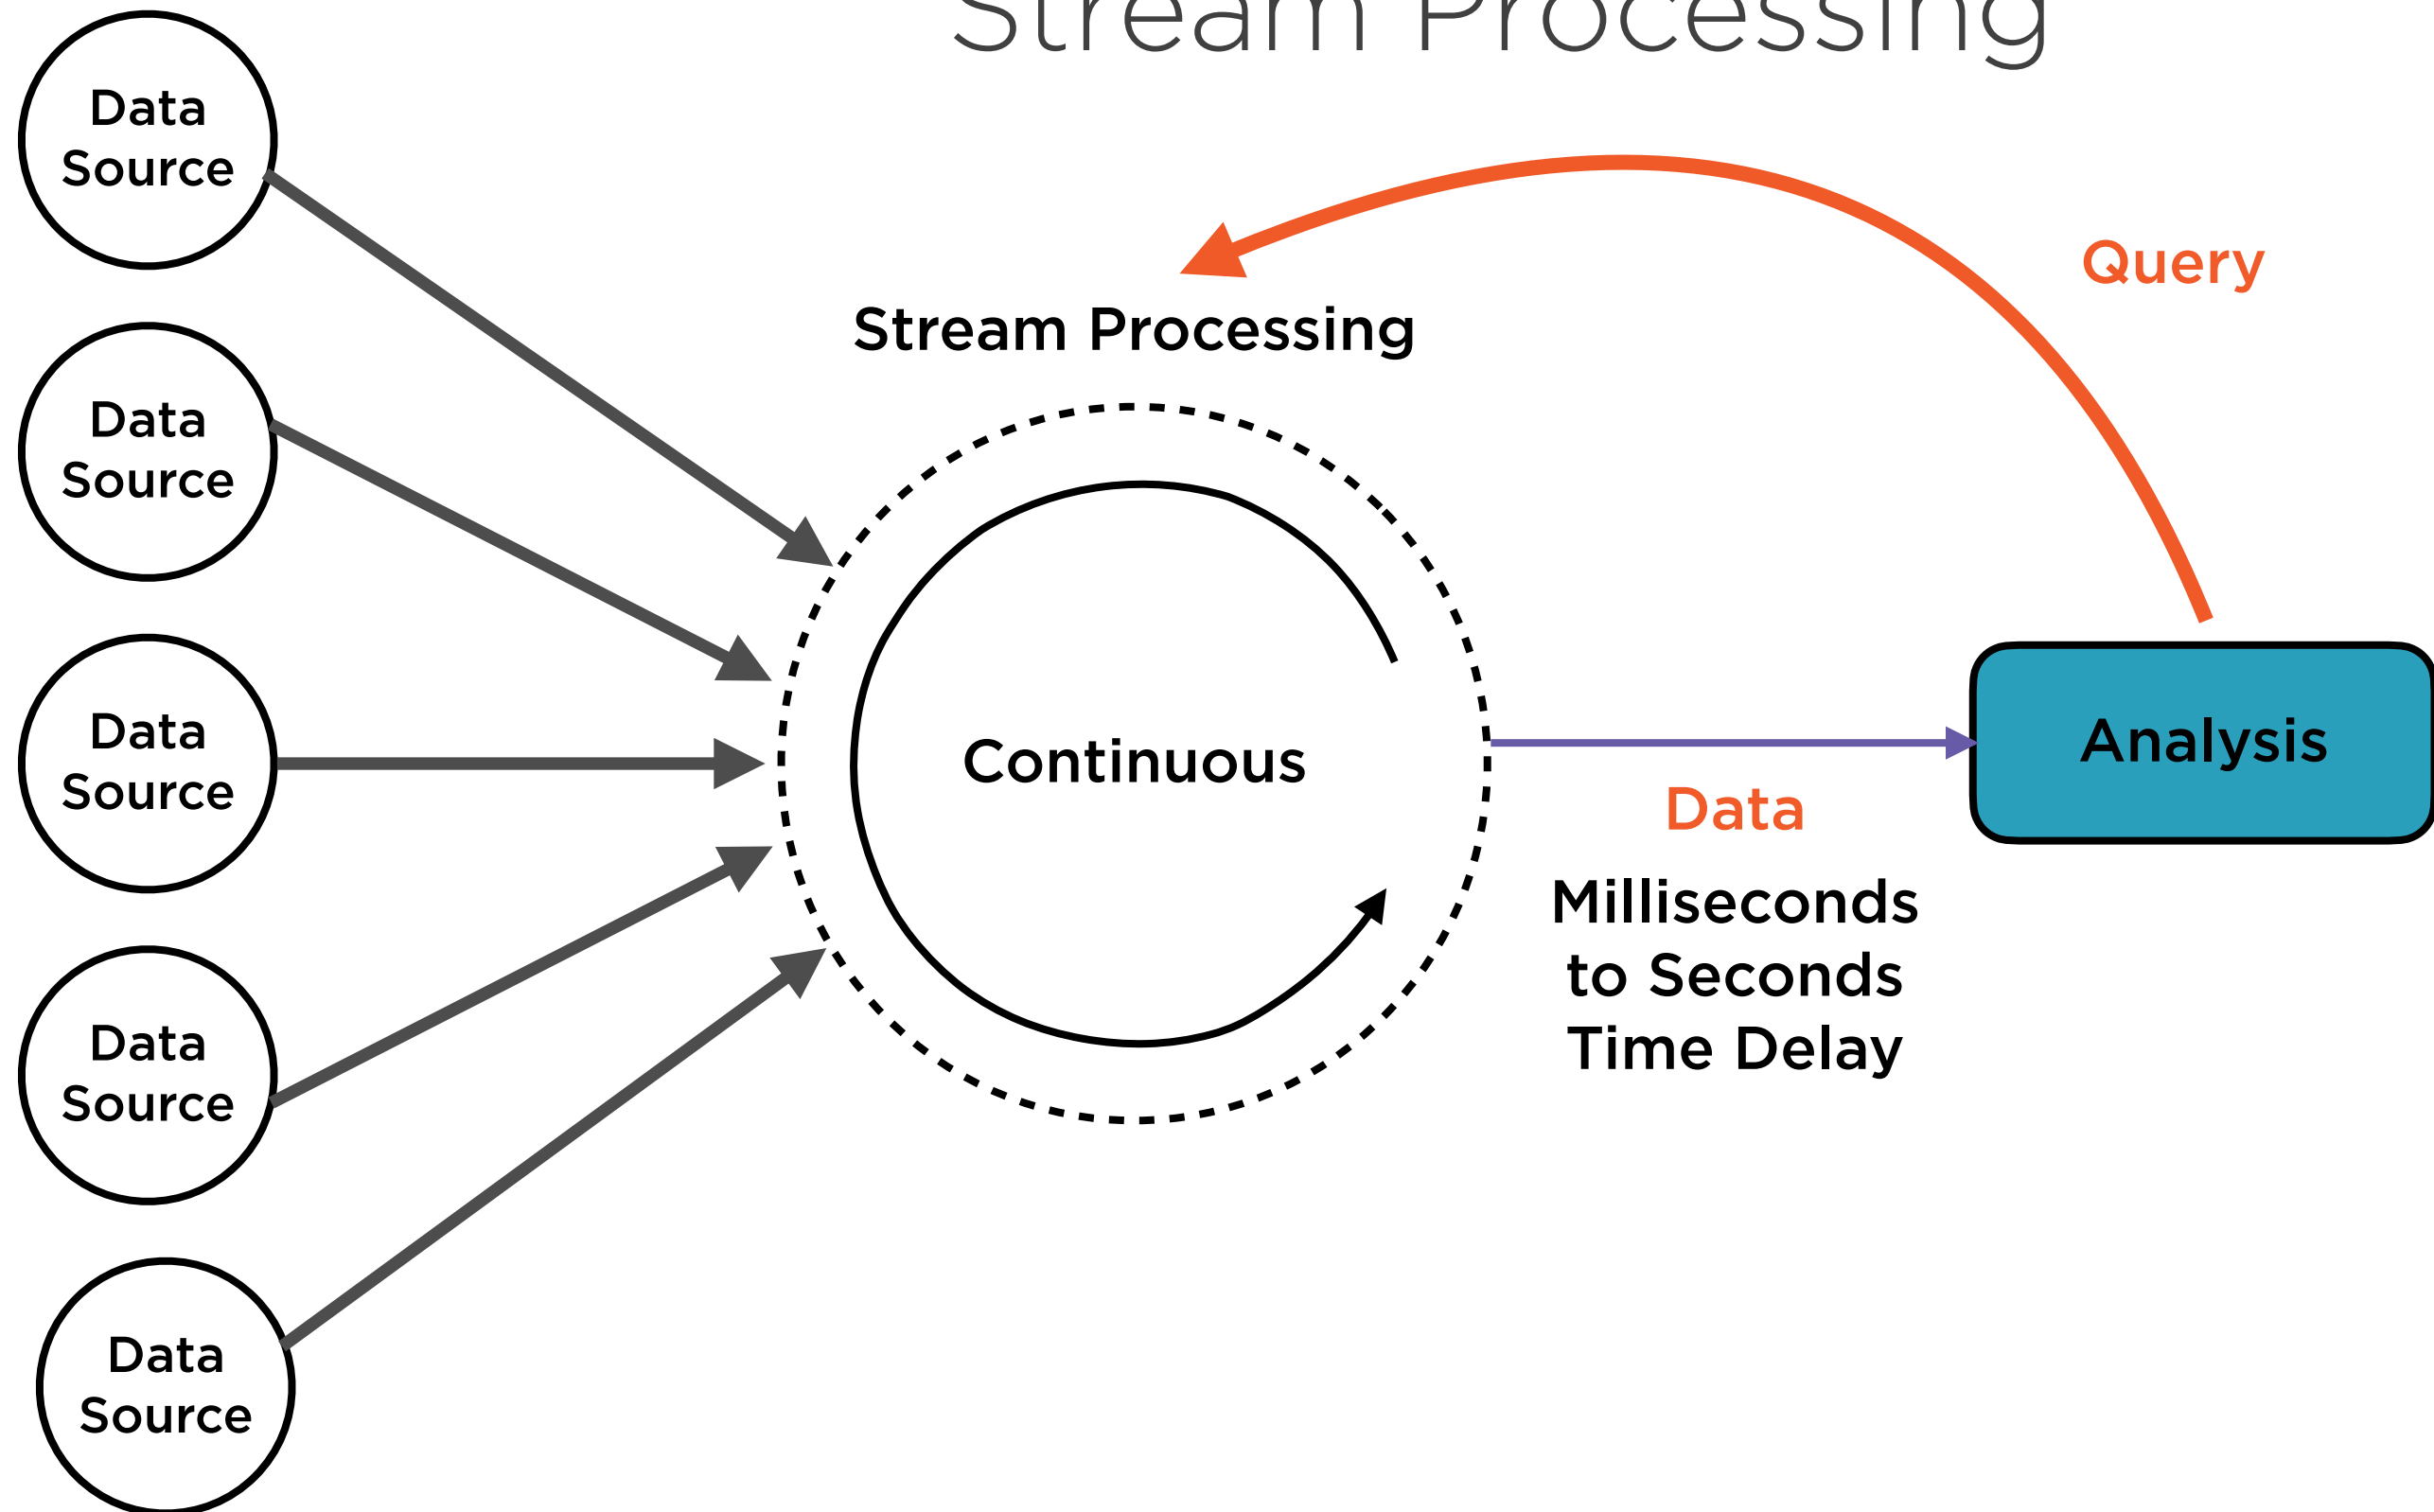


Input Data

**Input data is processed
with no time lag**



Stream Processing



Batch vs. Stream Processing

Batch

Bounded, finite datasets

Slow pipeline from data ingestion to analysis

Latency in minutes, hours considered acceptable

Periodic updates as jobs complete

Stream

Unbounded, infinite datasets

Processing immediate, as data is received

Latency usually must be in seconds, milliseconds

Continuous updates as jobs run constantly

Batch vs. Stream Processing

Batch

**Order of data received
unimportant**

**Single global state of the world
at any point in time**

Processing code “knows” all data

Stream

**Order important, out of order
arrival tracked**

**No global state, only history of
events received**

**Processing code does not know
what lies ahead**

Batch vs. Stream Processing

Batch

Payroll processing

No latency threshold

All employee data available
before processing begins

Stream


Fraud detection

Latency important

New data keeps coming - need
to detect fraud quickly without
slowing down legitimate
transactions

Stream Processing


Stream Processing



Data is received as a stream

- Log messages
- Tweets
- Climate sensor data


Stream Processing



Process the data one entity at a time

- Filter error messages
- Find references to the latest movies
- Track weather patterns


Stream Processing



Store, display, act on filtered messages


- Trigger an alert
- Show trending graphs
- Warn of sudden squalls

Stream Processing



Streaming data

Stream Processing



Traditional Systems



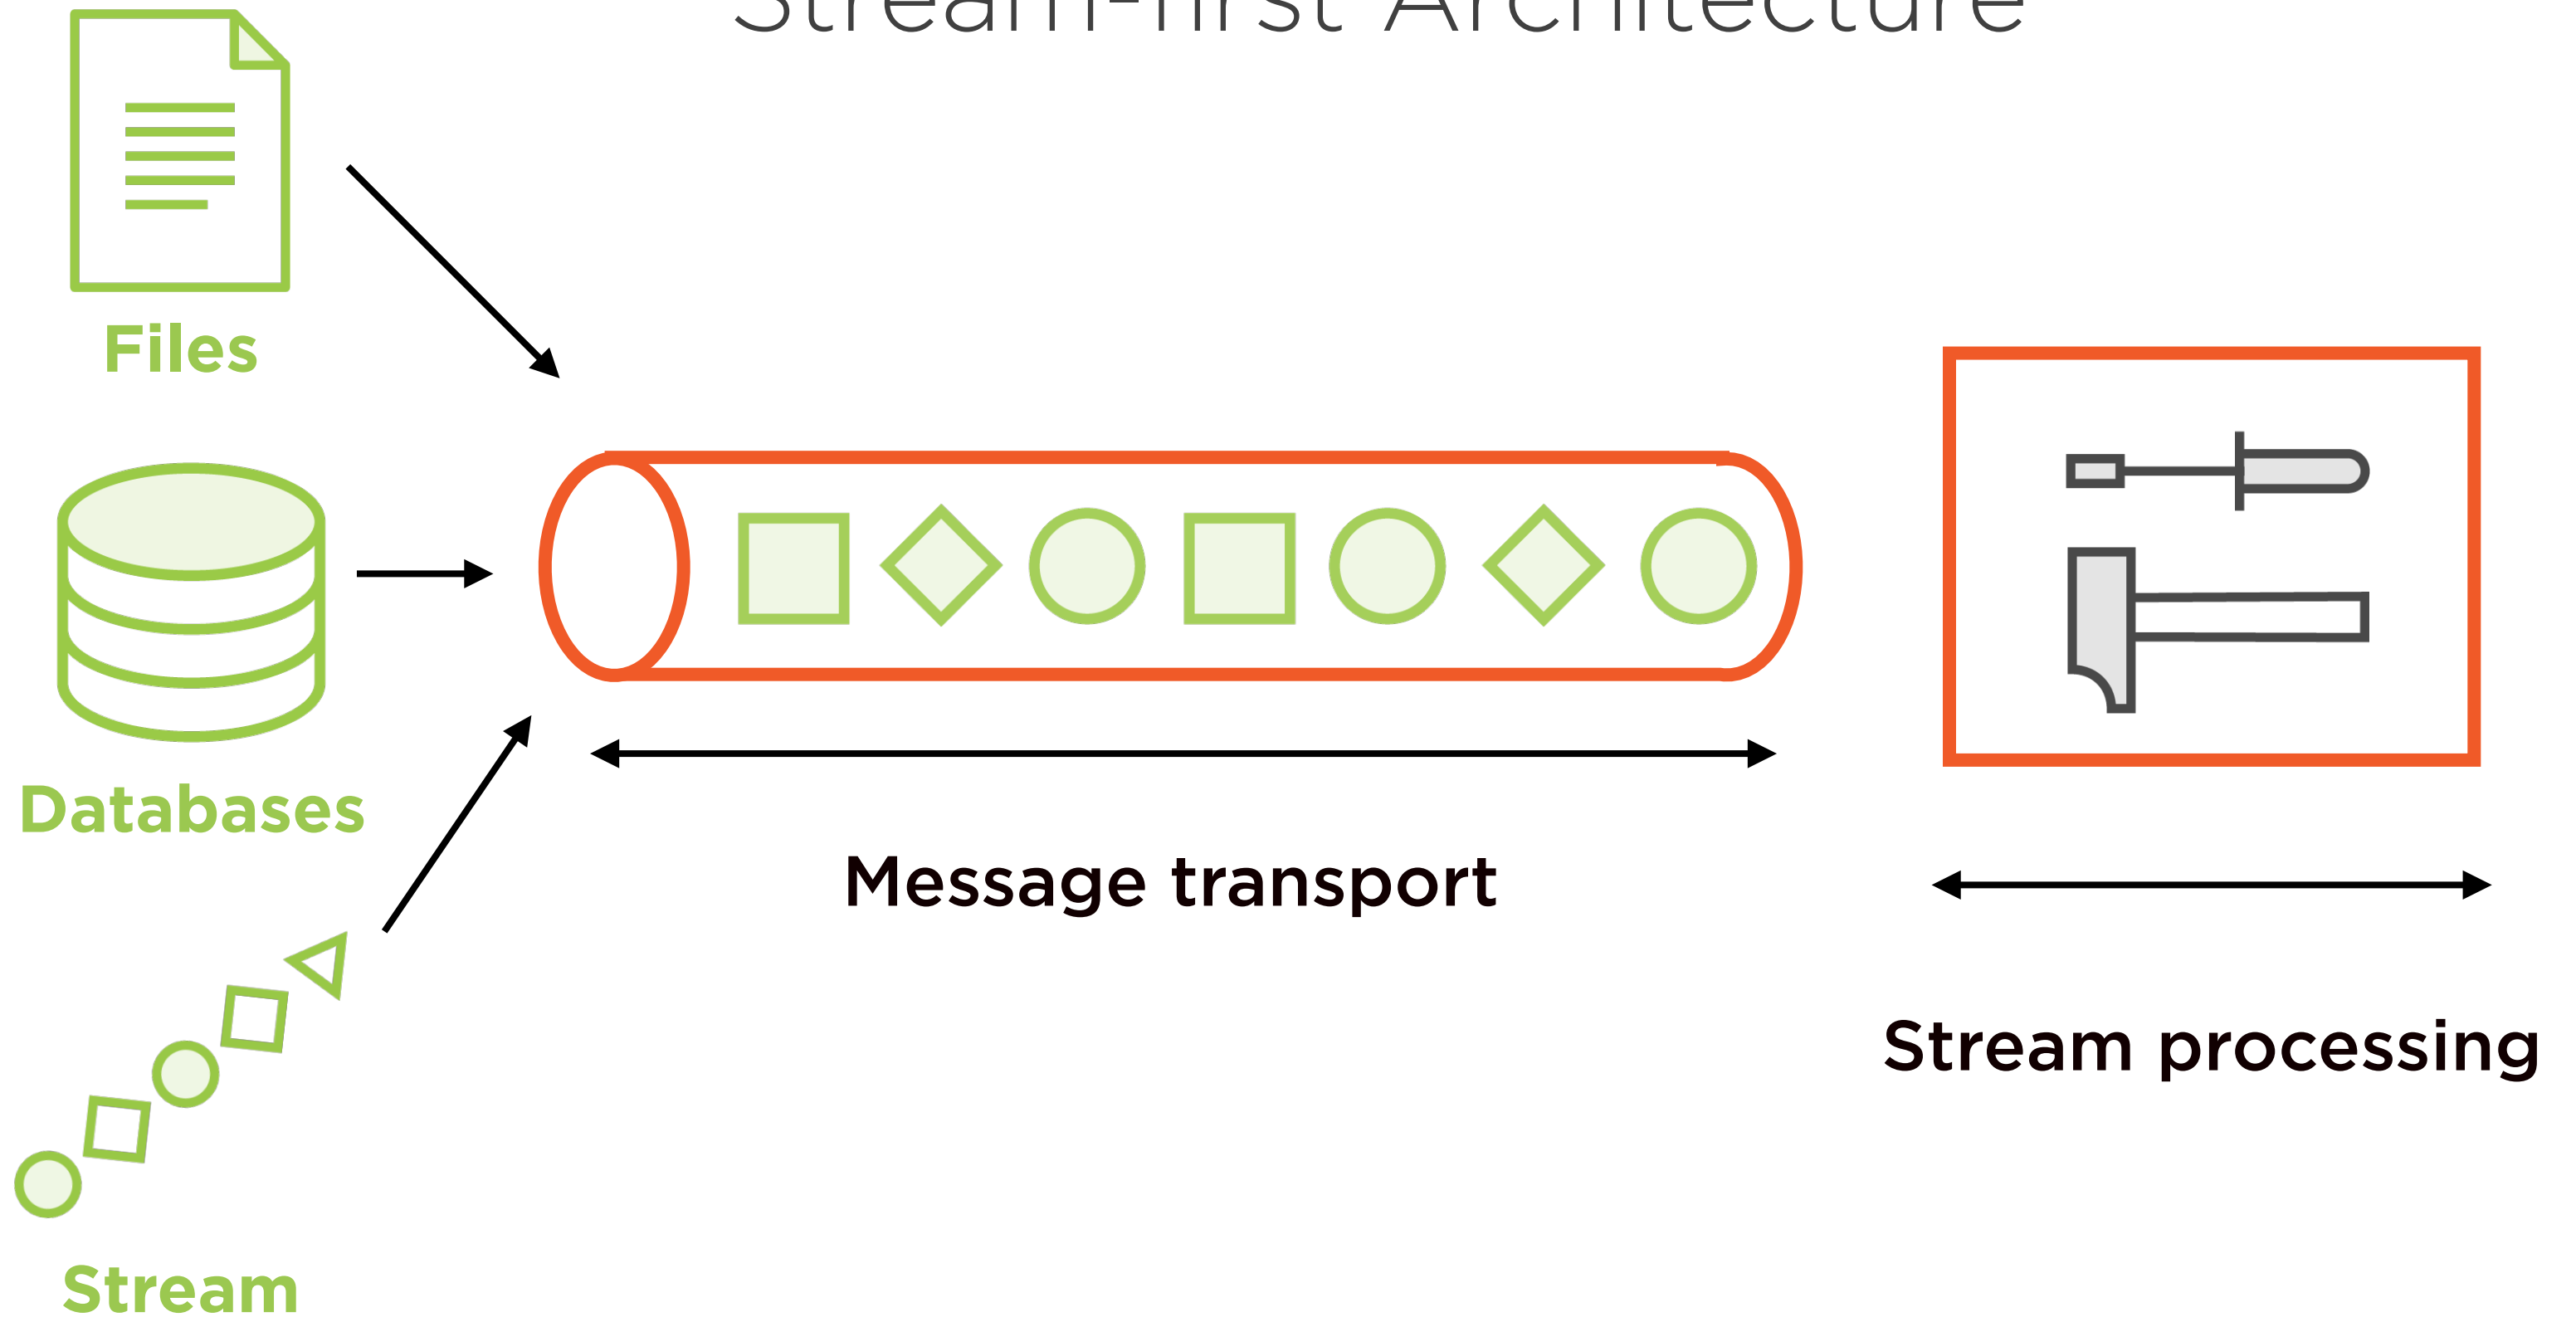
Files



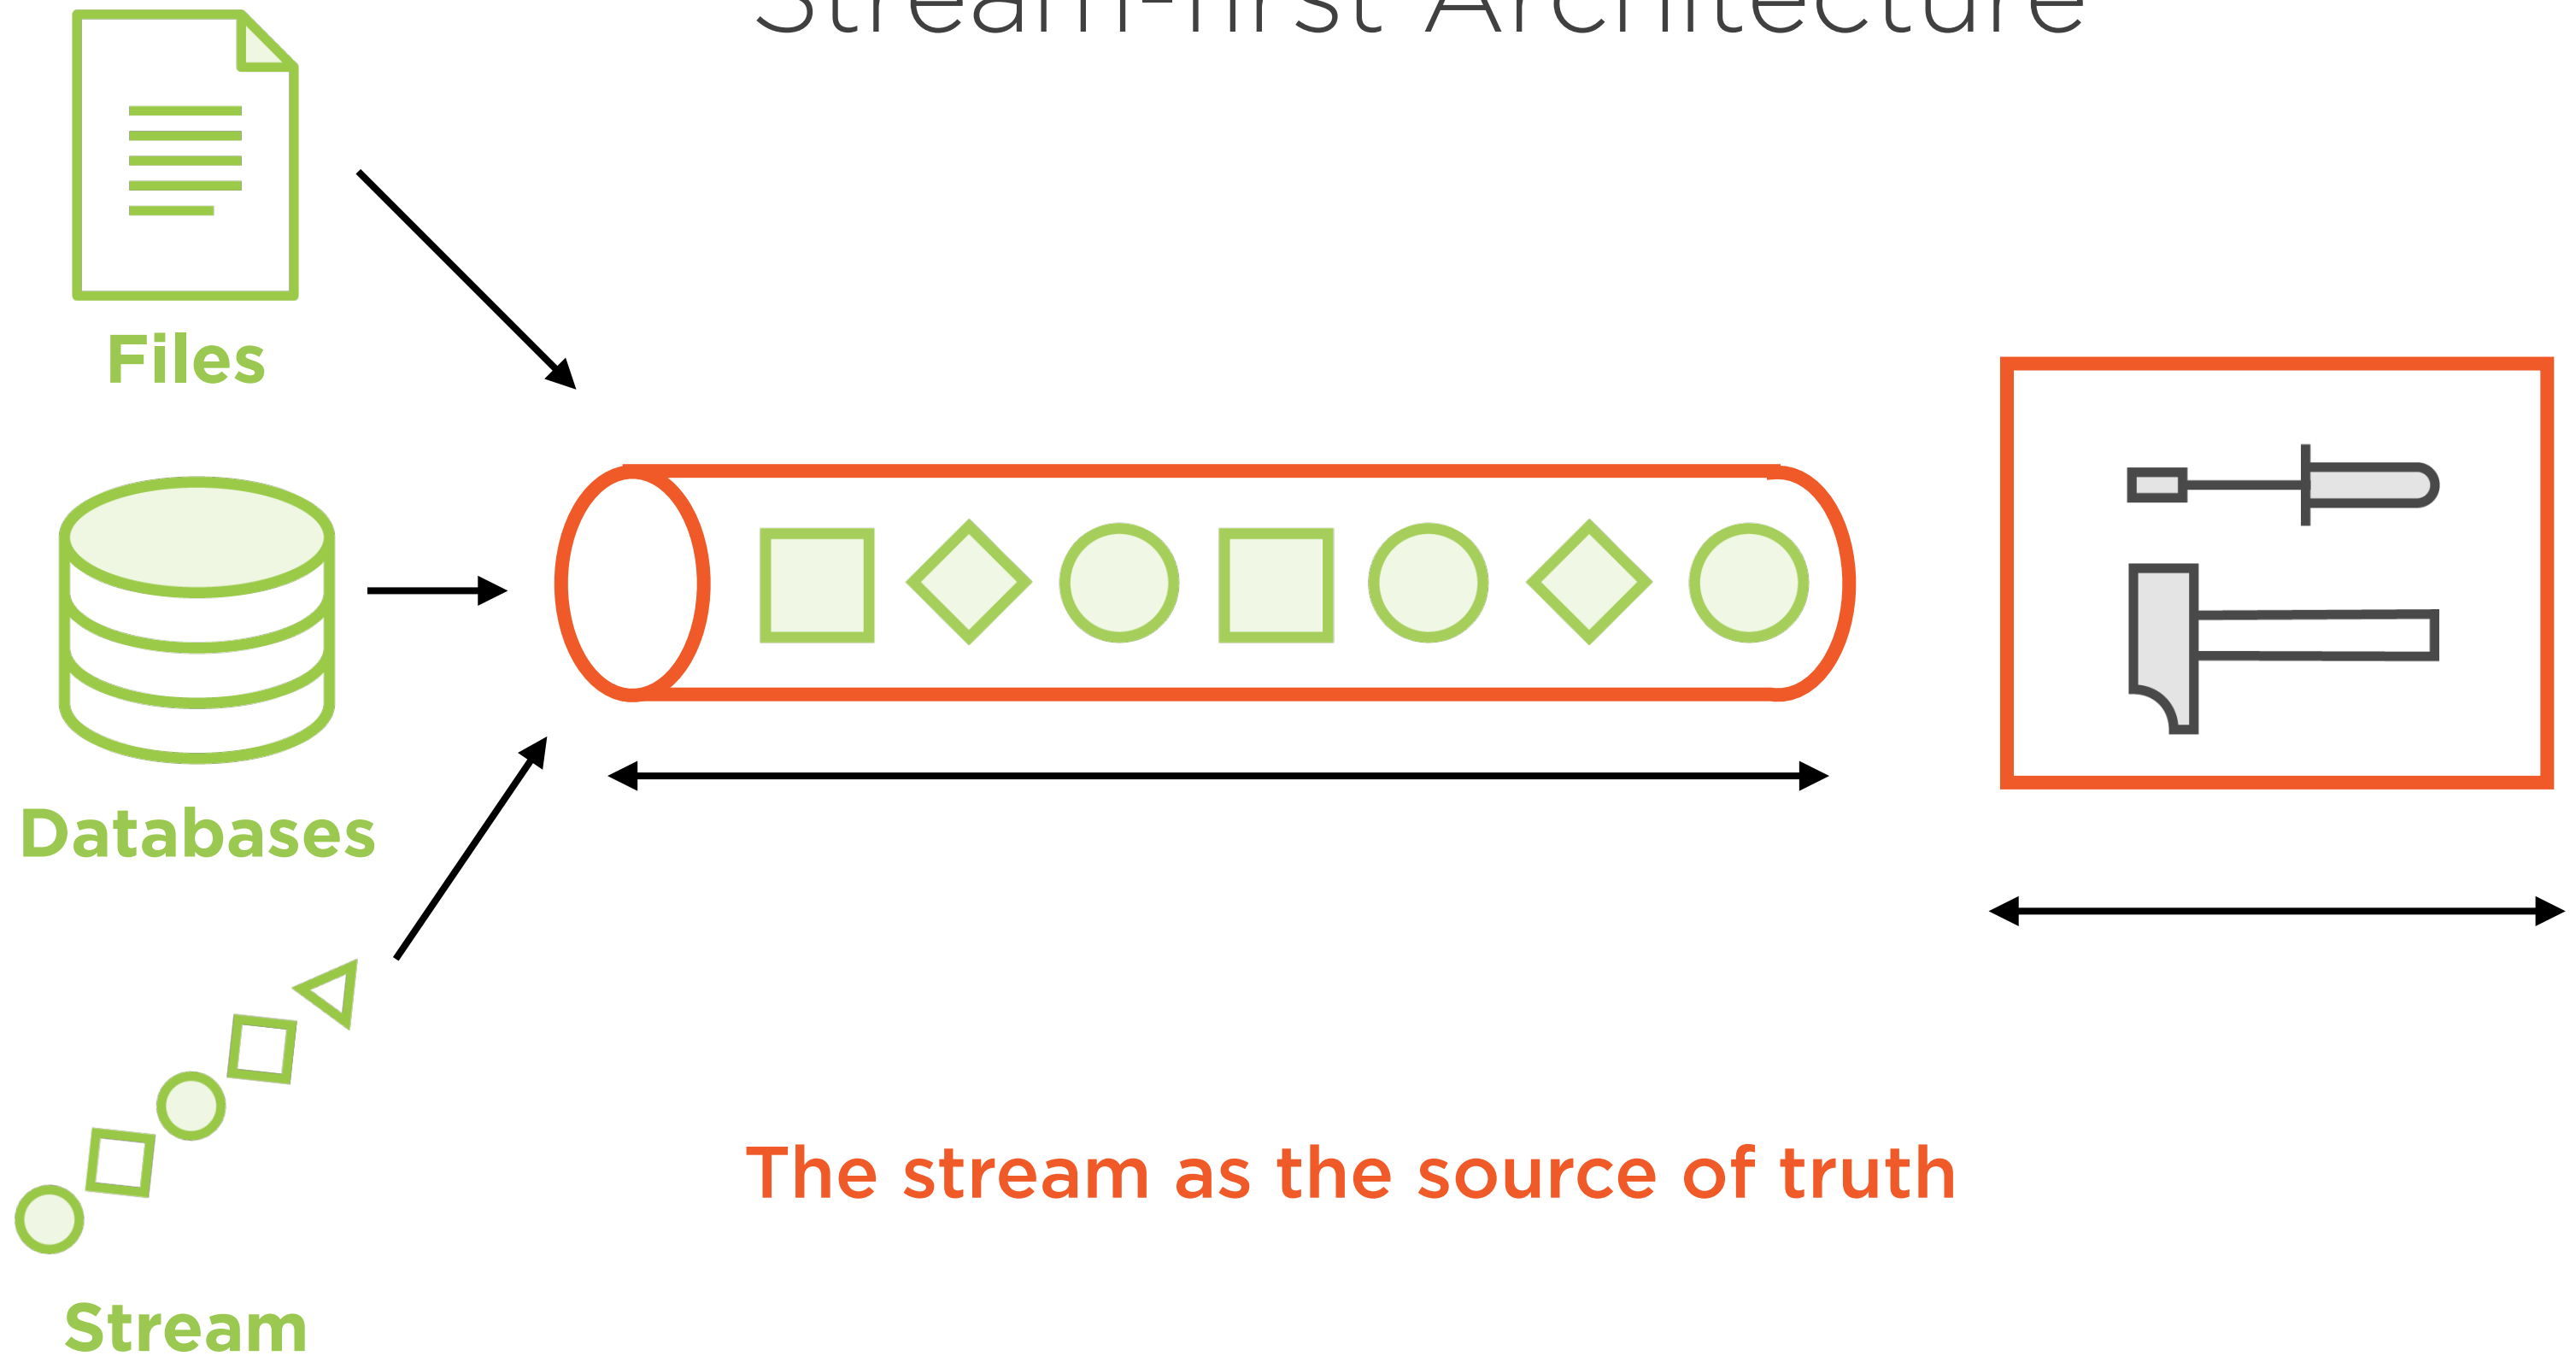
Databases

Reliable storage as the source of truth

Stream-first Architecture



Stream-first Architecture




Message Transport

Buffer for event data

**Performant and
persistent**

**Decoupling multiple
sources from processing**

Kafka, MapR streams




Stream-first Architecture




Files




Databases



Stream



Message transport



Stream processing

Stream Processing


**High throughput, low
latency**

**Fault tolerance with low
overhead**


**Manage out of order
events**

**Easy to use,
maintainable**

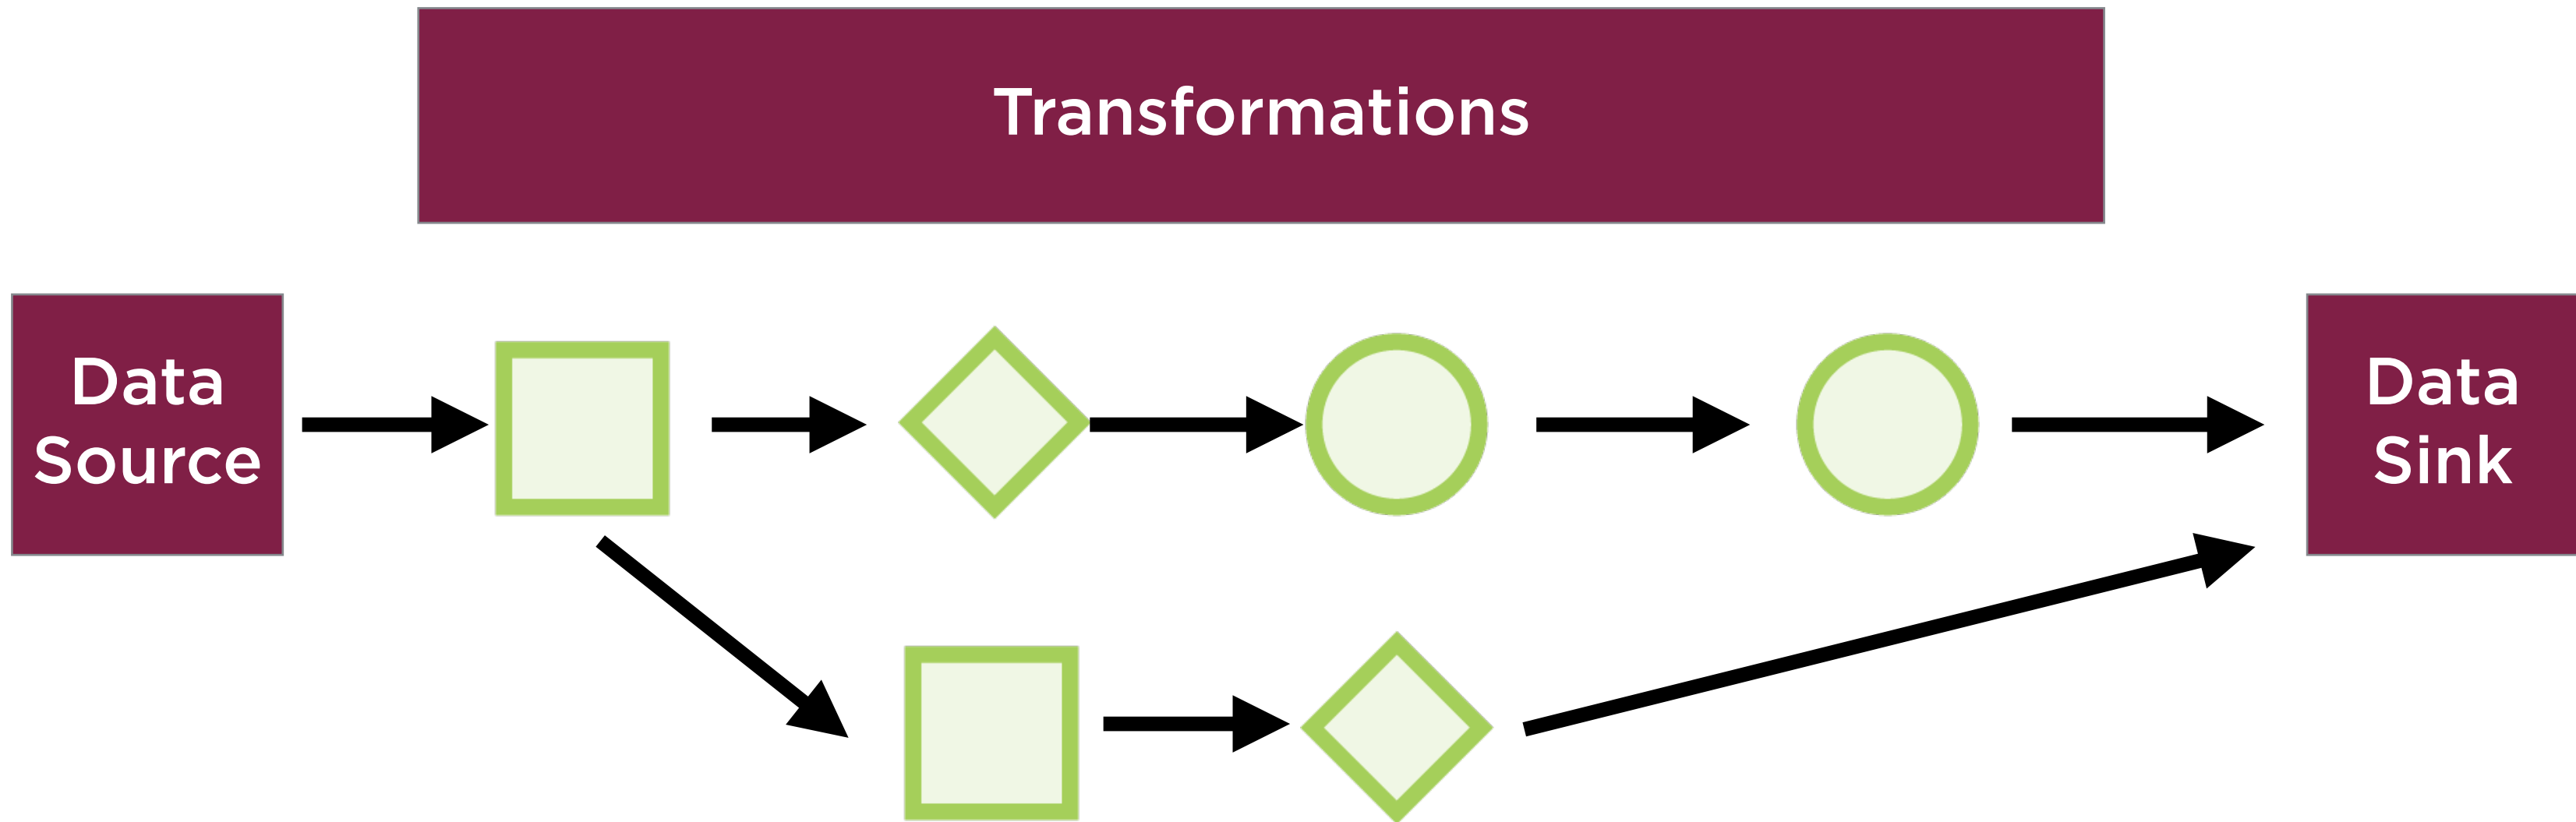
Replay streams



Stream Processing Model




Stream Processing Model



Transformations

A directed-acyclic graph



Stream Processing Models

Stream Processing Models



Stream Processing Models



Stream processing does not necessarily mean
continuous real-time processing

Micro-batch Processing




Run transformations on smaller accumulations of data

Collect say less than one minute of data


Process this micro-batch in near real-time

Micro-batch Processing




A stream of integers

Micro-batch Processing



Grouped into batches


Micro-batch Processing



If the batches are small enough...

Close to real-time processing

Micro-batch Processing



Exactly once semantics

Replay micro-batches

Latency-throughput trade-off based on batch sizes

Batch Processing for Streams



Latency, freshness of data are not considerations

Complex analytical operations

Joins on relational data

- Data might be in a data warehouse, need not be in an RDBMS

Continuous Stream Processing for Streams



Latency and freshness of data are **most important** considerations

Rate of arrival is high

- Latency in seconds/milliseconds only possible with continuous processing

Micro-batch Processing for Streams



Latency and freshness of data are important

but

Real-time processing is overkill

Rate of arrival is low/moderate

- Latency in seconds/milliseconds less important
- Acceptable latency possible with micro-batches

Stream Processing Architectures

Stream Processing Architectures

**Distinct Batch Layer
and Stream Layer**

**Unified Batch and
Stream Layers**



**The difference between these
architectures depend on how you treat
batch as well as stream data**

Lambda Architecture

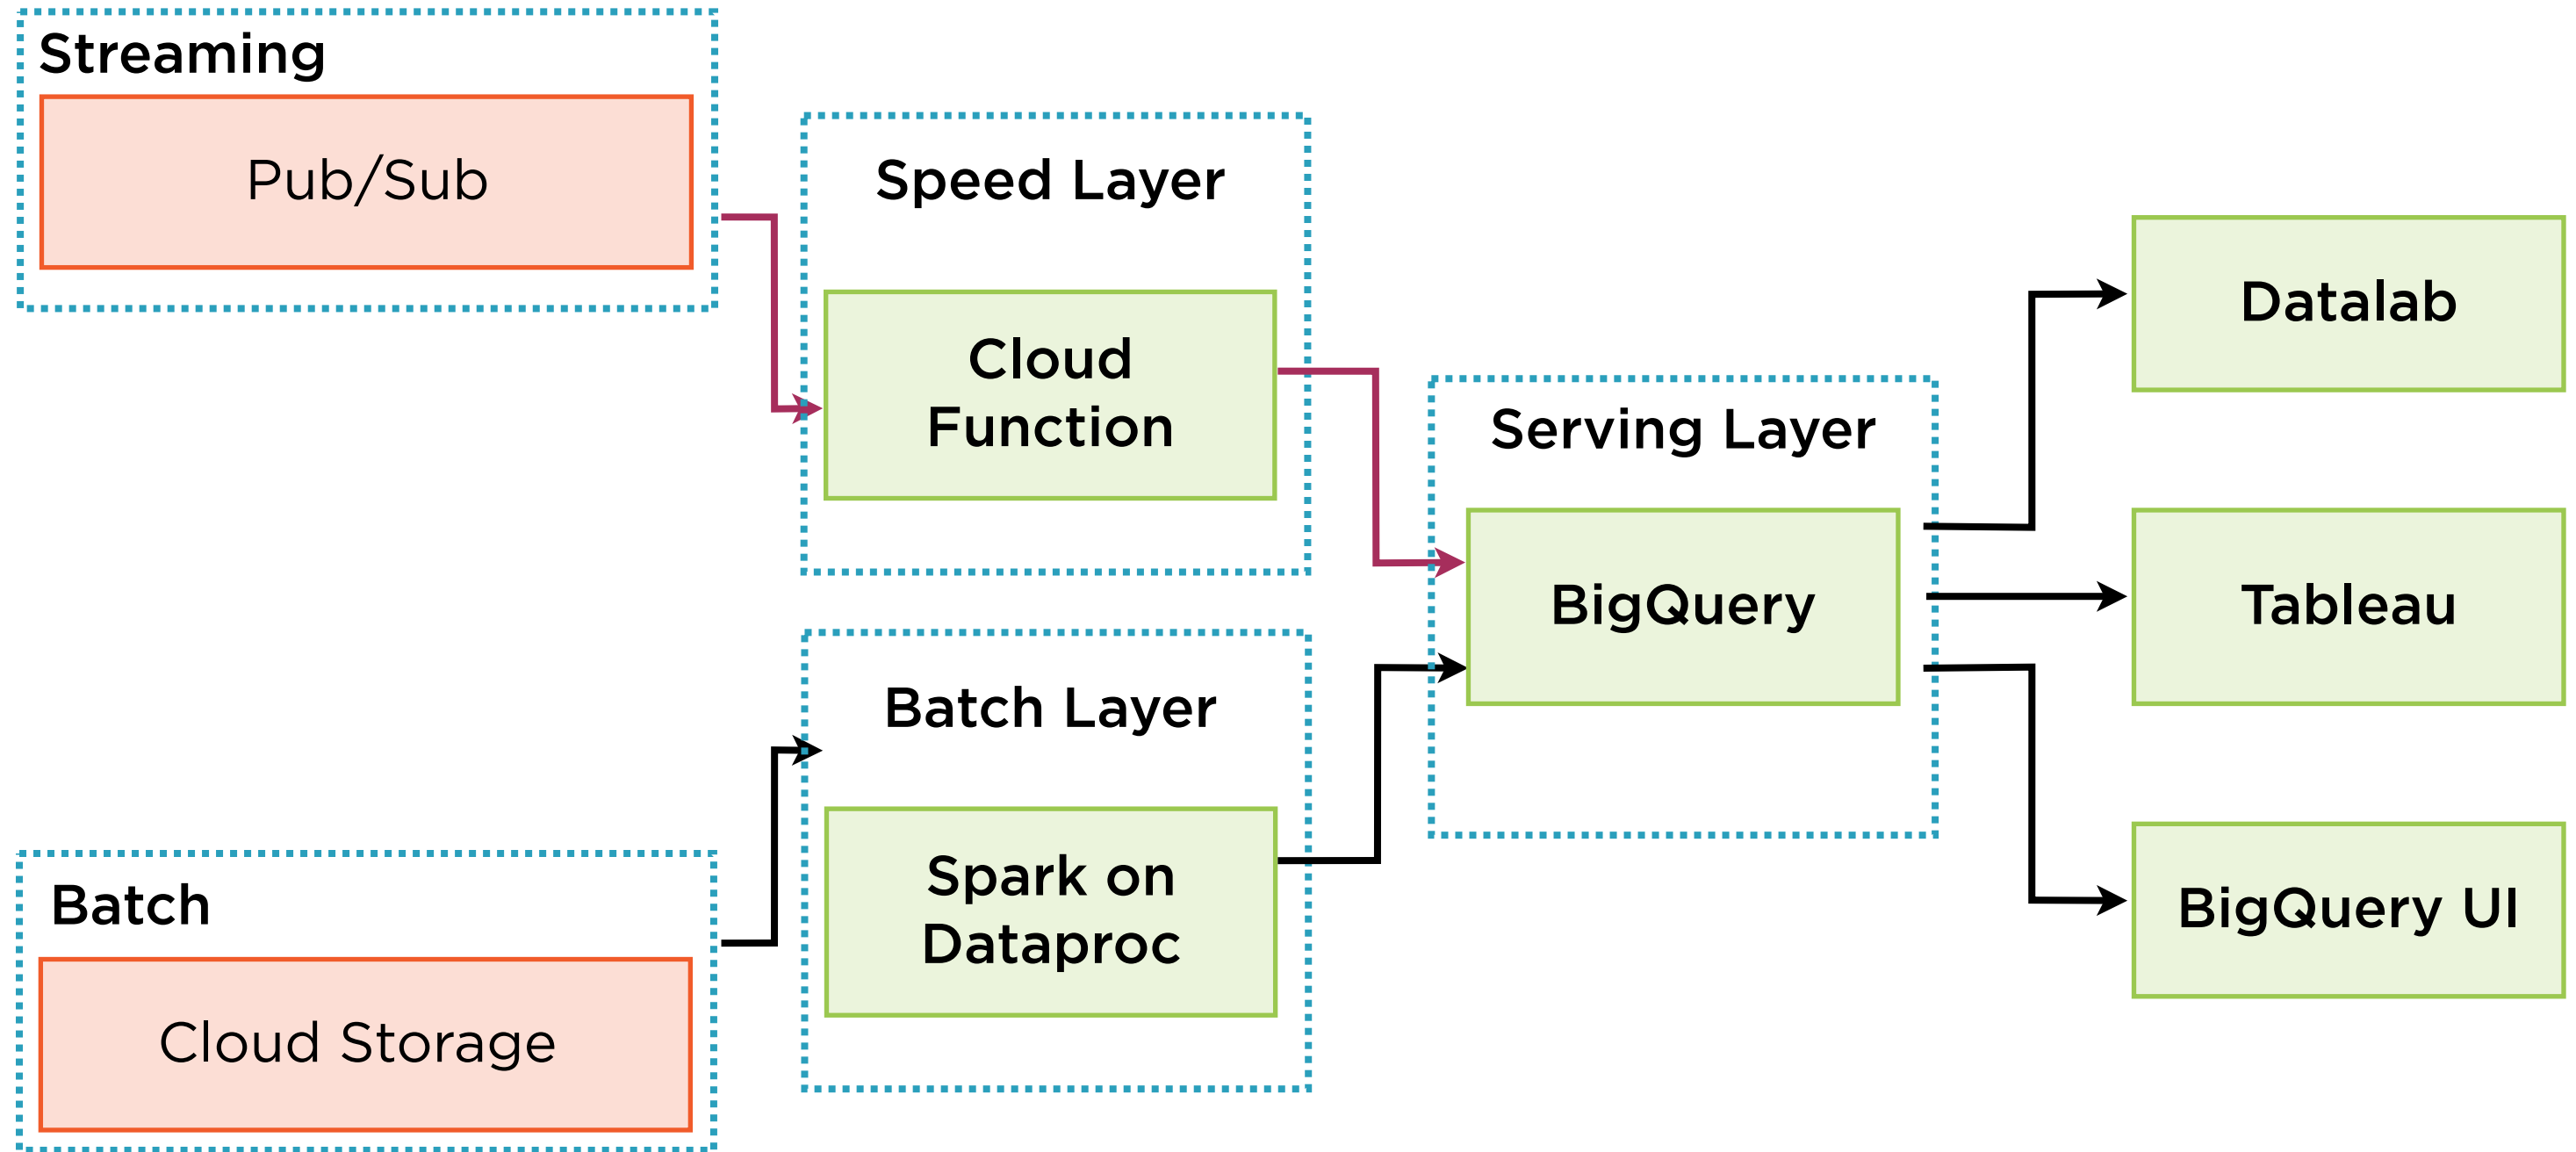


Run a streaming system along with a batch system

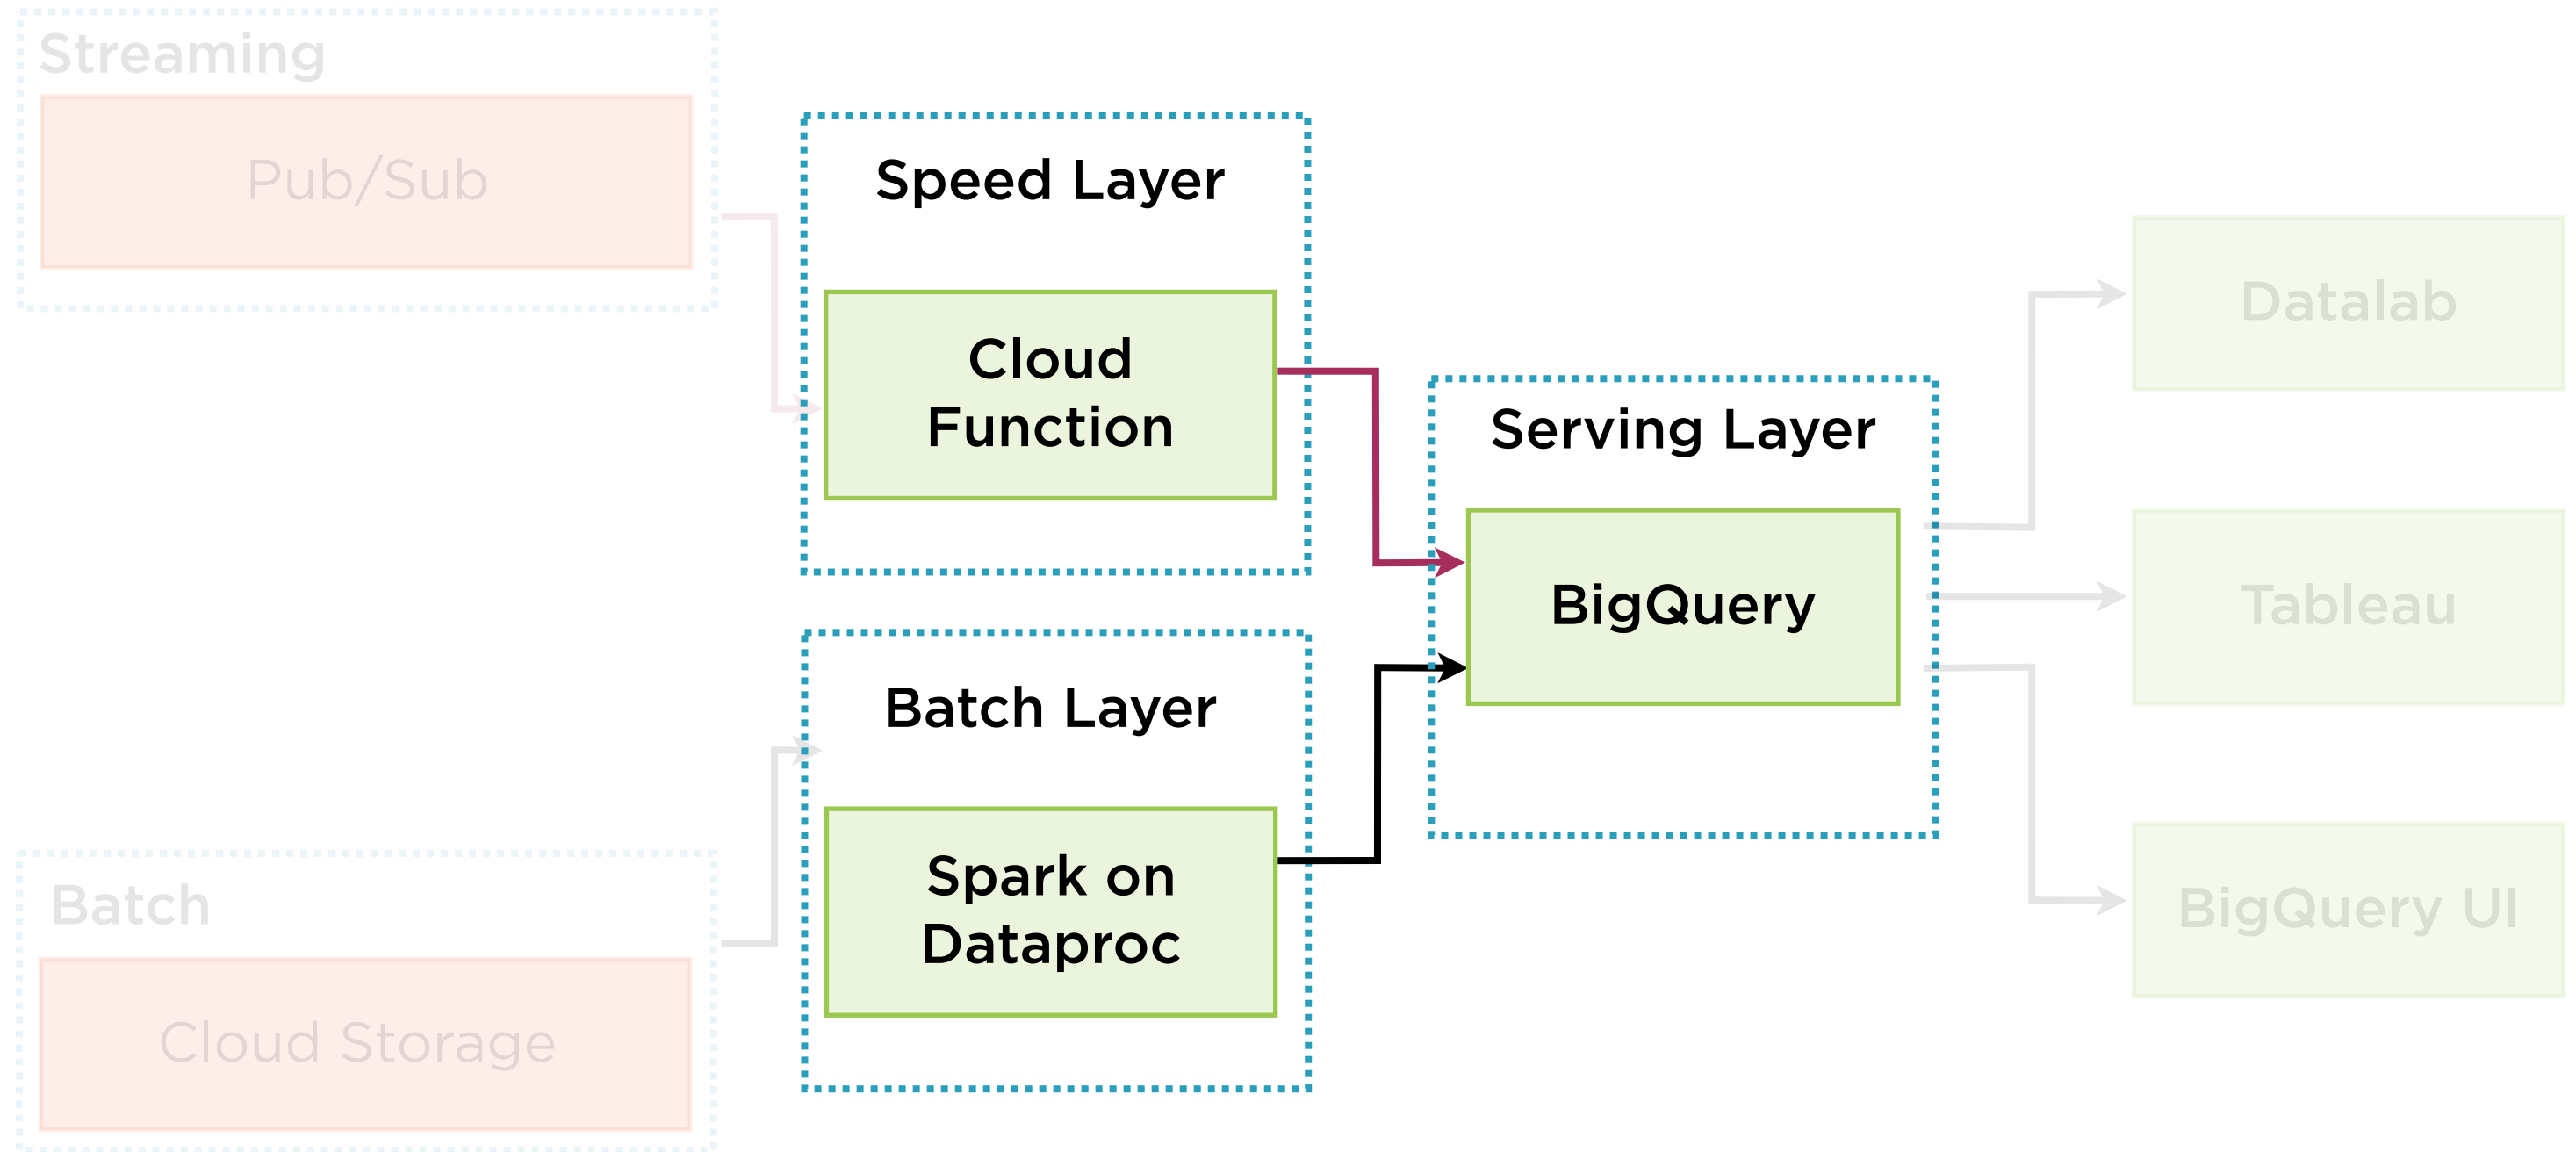
Streaming systems gives you low-latency but approximate results

Batch system ensures correctness, with higher-latency

Lambda Architecture



Lambda Architecture



Hybrid approach to batch and near real-time processing

Why Lambda?



Frameworks make separate batch and stream architecture choices

Because stream-first architecture offer poor performance for pure batch data

Optimizations for batch data are bolted-on, rather than built-in

Problems with Lambda Architectures




Batch and Stream layers perform the same computation, twice

- Batch computation is perfectly correct, but high latency
- Stream computation is low-latency, but often only approximately correct

Can lead to serious issues

- e.g. Training-serving skew in ML

Kappa Architecture



Kappa architectures tightly integrate batch and streaming


No separate code

Batch code simply fed through the streaming layer

In theory, eliminate training-serving skew

In practice, fragile and can be needlessly complex

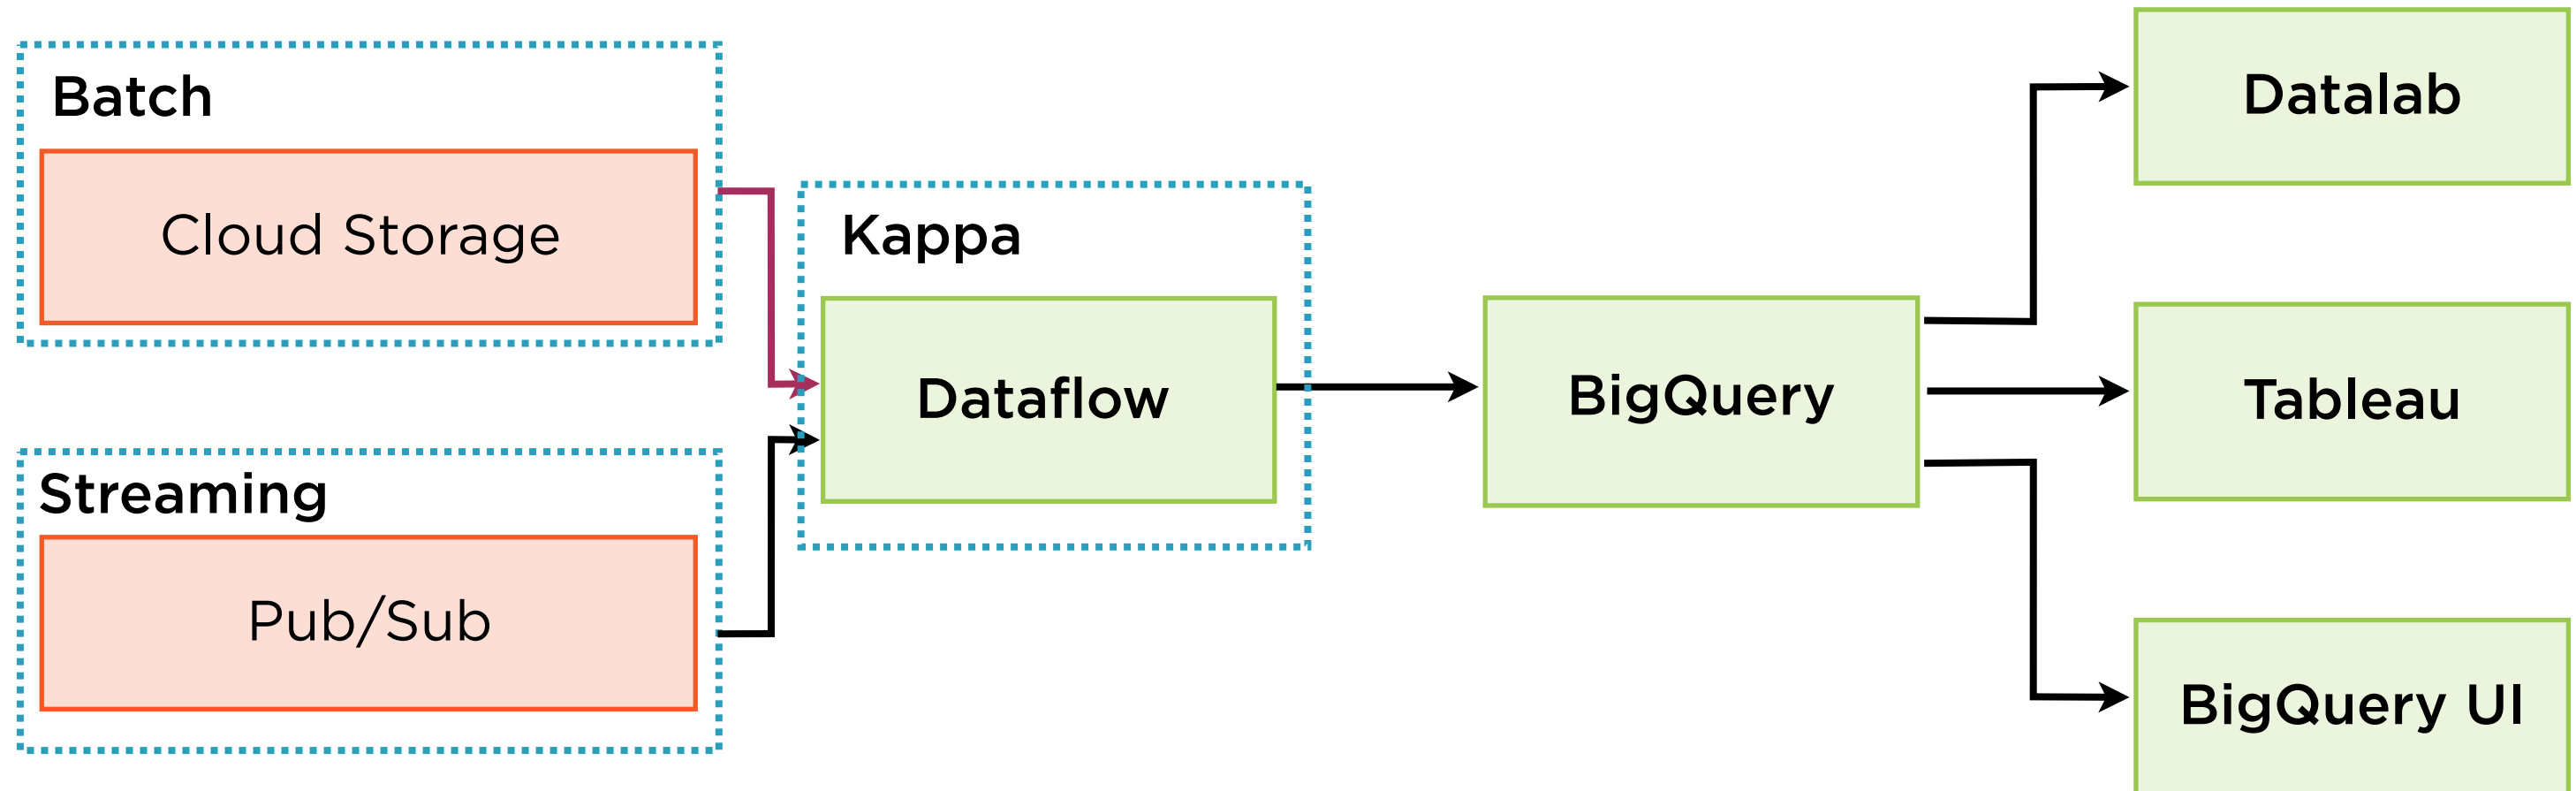
Kappa Architecture



“Batch is a special case of stream”

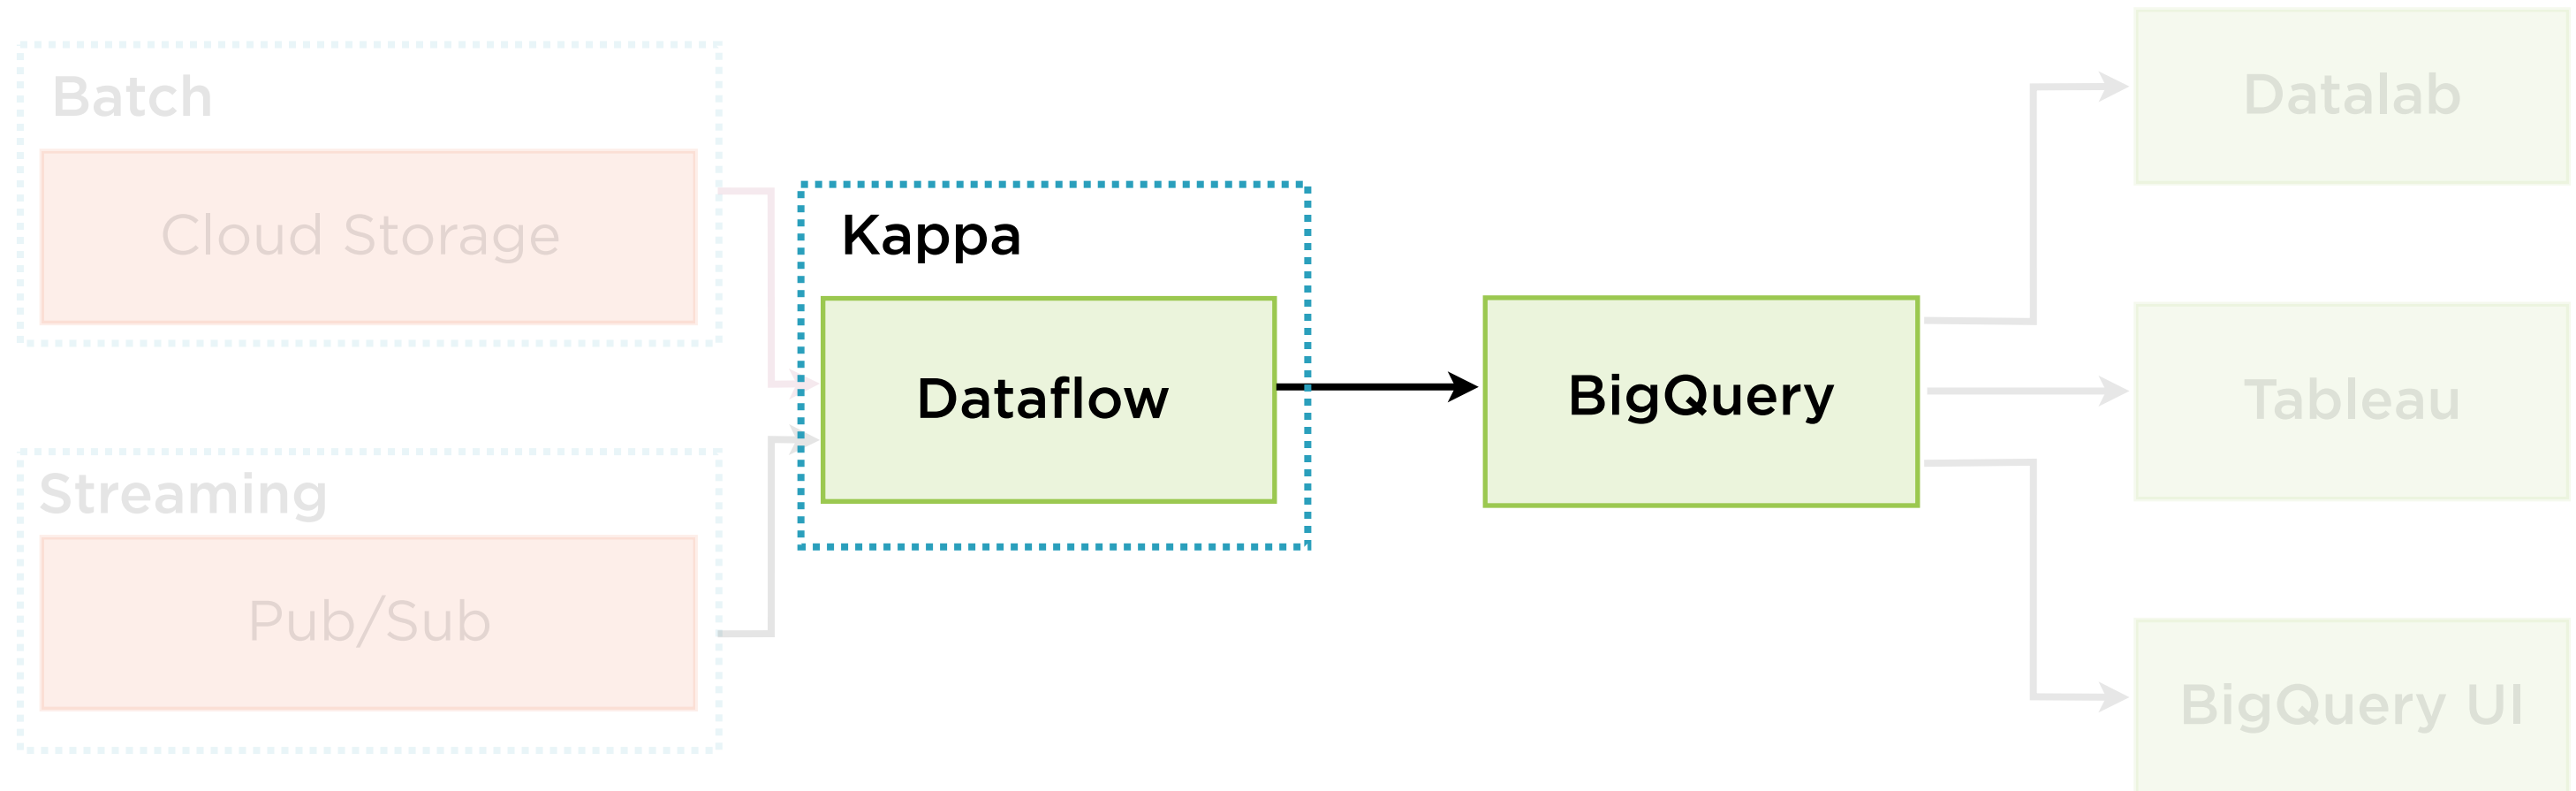
Well-designed streaming systems offer a superset of batch functionality

Kappa Architecture



**Process batch and streaming data using
the same code**

Kappa Architecture



**Process batch and streaming data using
the same code**

Stream Processing APIs

Distinct Batch Layer
and Stream Layer

Unified Batch and
Stream Layers



Unified Batch and Stream APIs are becoming more popular - but a Unified API can still rely on any of the architectures under the hood

Challenges of Stream Processing

Challenges of Stream Processing



Latency bounds

Dealing with late, out-of-order data

Guaranteeing reliability

- “Exactly-once, ordered processing”

Security

- Encryption
- Authentication
- MITM attacks

Challenges of Stream Processing



Dimensions of scaling

- Number of senders
- Number of receivers
- Number of messages
- Organizing messages (e.g. topics)
- Size of messages

Summary

Batch data and bounded datasets

Streaming data for unbounded datasets and real-time processing

Micro-batch processing and continuous processing

Lambda and Kappa architectures

Challenges in real-time stream processing

Up Next:

Introducing Apache Beam for
Stream Processing
